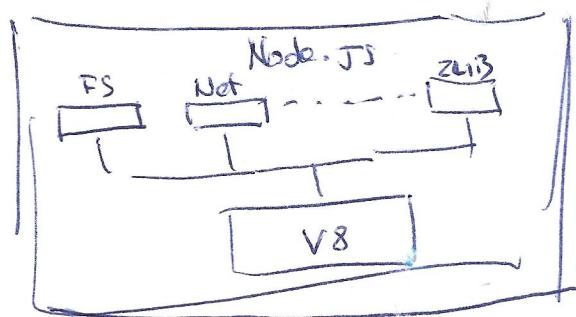


Node.js expose low level I/O \rightarrow FS - threads - Network



JavaScript Interpreter — is more like a JIT Compiler.

We are more interested in the fact that V8 provides a C++ API into the JavaScript that is executing. This API allows us to write C++ code that can be loaded by Node.js and this C++ code become fully callable from JavaScript by JavaScript in C++ and even

The API allows us to hold resources allocated by JavaScript in C++ and even

call JavaScript functions from C++.

Three main components

1. C++ that holds our C++ addon

2. Binding.gyp

3. JavaScript file that use the AddOn

Hello World
Example

Registering the module

Get a reference to the V8 execution context

JavaScript Memory Allocation

V8 JavaScript code Runtime — libuv Event loop implementation

Isolate is an independent instance of the V8 runtime -

Execution context | Memory management / gc | Heap | Pool of memory

It's not true that we want build native addons for speed. There are some tasks that perform well in JavaScript.

All that said, C++ should deliver at least a 2x run time improvement over JavaScript for most CPU-bound operations or tasks.

First reason CPU-heavy task to complete in a Node.js Application. You already have some C/C++ code that you want to integrate in your Node.js application.

Migrate over C++ code that is already important to your application or integrate it through native addons.

Performance yes but you pay attention to the cost of data copying.

We always copy data out of V8 before utilizing fully.

Node.js and V8 Engine

How is this done? What is providing this bridge between these two very different languages?

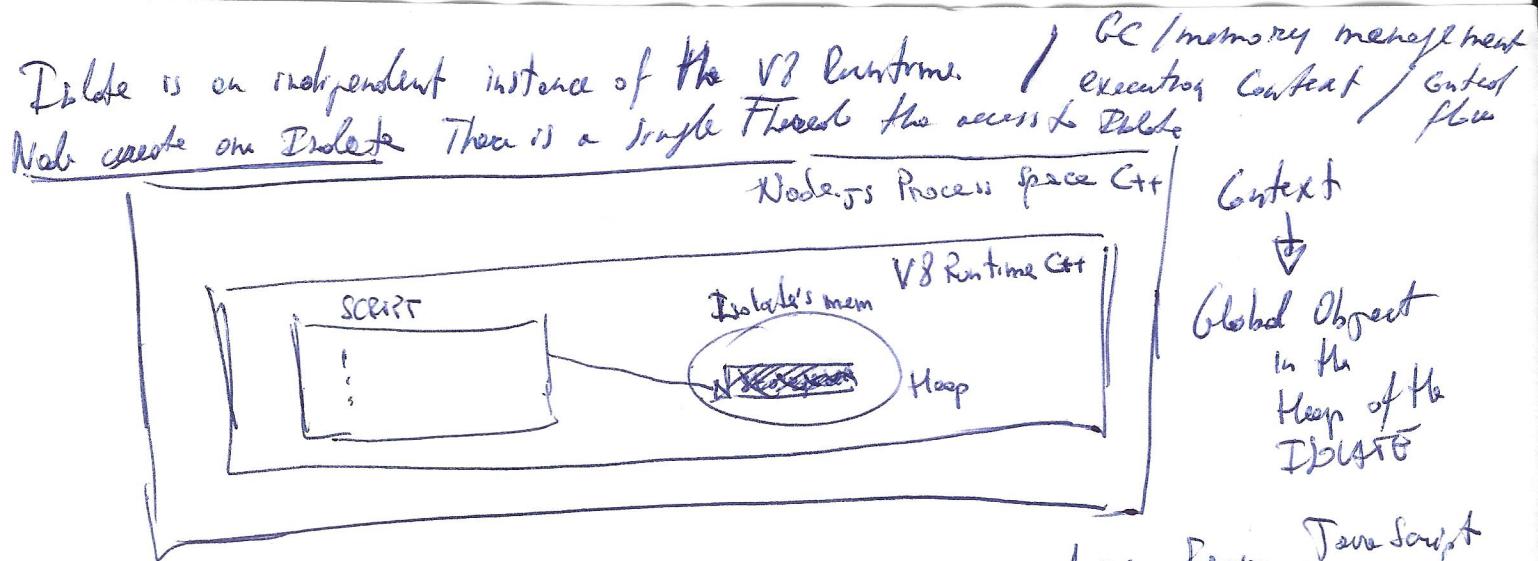
Google V8 JavaScript library



C++ Library which provides a high performance execution environment for JavaScript code.

Node.js is just the glue and support code around a few key components.

1. Google V8 JavaScript Engine
2. C++ modules for low-level I/O
3. Event loop



V8 not only needs to know the references originating from JavaScript
but also in the addons.

Handle → Object that represents the existence of such a reference in C++.

Local → Allocated by C++ call stack

That's to process to get ~~variables~~ exchange variables value between
JavaScript and C++ - Primitive and non Primitive Data Type.

Best Error handling

Asynchronous API with Worker Thread

Event Loop Thread

2 Threads

Worker thread managed
by
libuv

Executing JavaScript code
We stay in this thread
when we cross over in to
C++ addons

This is the thread that we
don't want to stall by doing
heavy calculations

Each thread has its own stack

Event Loop Thread

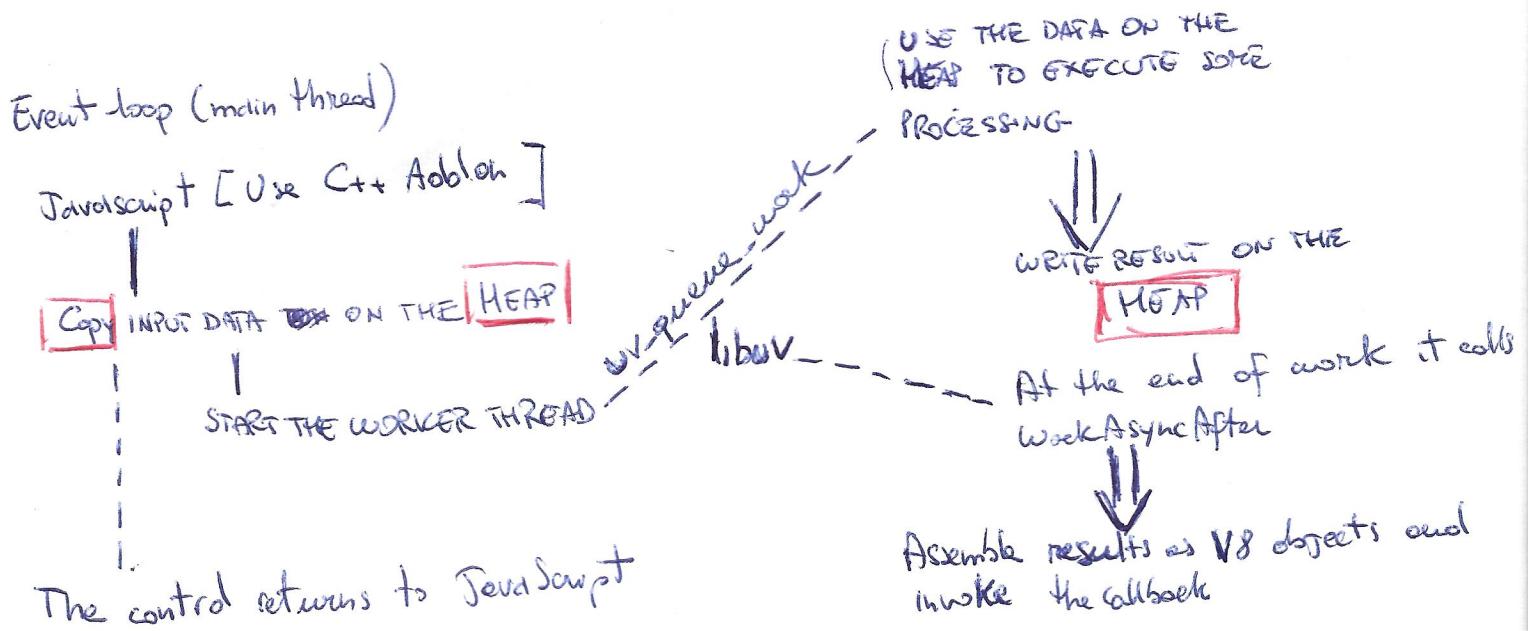
Worker Thread

It is not possible to share stack variables
between the event loop thread and worker threads
But these two threads share the Heap and we
are going to put our input and output data -

JavaScript is implicitly single threaded. V8 is built around the notion that state within JavaScript is strictly accessible by ~~one~~ one thread at a time.

YOU CAN'T ACCESS V8 MEMORY OUTSIDE THE EVENT-LOOP'S THREAD

To avoid the violation of this rule, you need to create a copy of the input/output data.



Worker DATA

```
struct WorkerData {  
    uv_work_t request;  
    Persistent<Function> callback; }  
It will be kept in the scope by V8 and  
not garbage collected.
```

Copying input and output data is being done in the event-loop if it takes long time we're blocking the event-loop - In that case we have two problems:

1. Copying data might be a waste of memory
2. Copying data might take long, which ties up the Node.js event-loop

~~ANSWER~~
NAN provides a set of memory ~~operations~~ and ~~subtleties~~ that achieve ~~too fast~~ goals:

Some code has been written, it must be compiled into the library `binding.gyp` → Describe the build configuration of the module using JSON-like format

This file is used by `node-gyp` a tool written specifically to compile Node.js Add-ons

`node-gyp configure` → Generate project build file for the current platform
`node-gyp build`

`build/Release` → `**.node`

Use `require` to import the module.

Direction of Node.js and V8 API

V8 API can and has, changed dramatically from one V8 release ~~to the next~~ to the next.

~~With code change~~ NAN provide a set of tools that addon developer are recommended to use to keep the compatibility between the past and the future.

API for building Native Add-ons independent from the underlying JavaScript Runtime (e.g. V8) This is API will be the Application Binary Interface (ABI) stable across different version of Node.js.

N-API is independent from the underlying JavaScript runtime (ee V8)

N-API calls ^{Return} STATUS CODE \Rightarrow napi-status
This return value is passed via out parameter

napi-value — ABSTRACT ALL JS VALUE

In case of error additional info can be obtained calling napi-get-last-error-info

What is provoking this war bridge between these two very different languages?

Google V8 JavaScript Library

V8 OpenSource  JavaScript execution engine. V8 is a C++ library which provides a high performance execution environment for JavaScript code.

V8 provides C++ API into JavaScript that is executing.

↓
C++ code can be loaded from ~~FS~~ by Node.js
and this code become fully callable from ~~Node.js~~
your JavaScript code.

To hold resources allocated by JavaScript in Cell and even call Javascript function.

Initialization code It means that bootstraps code that will register a module to the system as you forget. When the module is registered it ~~assure that~~ ensures to call Init method

Passing parameters from JavaScript to C++
Inspecting an argument before using the them.

It's possible to shield yourself from ABI changes using Native Abstraction for Node.js

Java Script \rightarrow C++ Poss. like as a single copy

1) JS/VS \rightarrow C++

2) Execute the processing / calculations

3) C++ \rightarrow JS/VS

Any existing code
C++ can be used
without much modification.

Copying data to a
timer

Integration Pattern: Data Transfer

We are doubling memory consumption and its also costing us processing time to do all this.

The event loop implemented by a library called libuv should never stall on I/O or heavy CPU because, in theory, there is likely other work to be done.

Q) You are C++ to perform long running task & where C++ will run from JavaScript

Keeping those points in mind, the better and unlockable approach fails miserably. The worker thread can only access V8 objects when the event loop is sleeping.

Where is that date coming from?

Is it coming from a database? Retrieve it in C++ not in JS
Is it coming from a file? Read it from C++

ObjectWraps

If your code doesn't sit somewhere already - JavaScript code is actually building it incrementally over time, the another option is to store the date in the C++ addon directly. As you build state instead of creating JavaScript arrays, objects, make calls into the C++ addons to copy state incrementally into C++ state structure that will remain in the scope.

ObjectWrap → Use C++ Object directly from JavaScript
Buffer Object → This technique works because Buffer object allocates state outside of V8 so it can be accessible from the threads.

Object Wrapping

There is a way to create the bridge between JavaScript and your add-on but you have to create a new kind of ~~present~~ object

`node::ObjectWrap`

It contains all the right tools to allow us to connect JS code with

C++ code

Classes that extend `ObjectWrap` can be instantiated from JavaScript using the `new` operator and their methods can be directly invoked from JavaScript

The `new` word refers to the way to group methods and state. There isn't nothing magically.

Neither constructor ~~nor~~ destructor is ~~public~~ public. They will not be called directly from outside (much less from JavaScript)

But has the responsibility to ~~call~~ call all class functions that will be called from JavaScript for the ~~extern~~ object.

But is called by `InitPoly` which is the entry point to our add-on and it's called as ~~as~~ soon as our add-on is required.

The new C++ object then gets some `NodeJS` sprinkled into it by calling

its `Wrap` method

Handle the constructor because we can create functions

object with the `new` operator and simply calling functions

Node.setPrototypeMethod to add a prototype method on object
SetInternalFieldCount instructs V8 to allocate internal storage slots
for every instance created using template.

It's not Node.js that defines the API for which we interoperate
with JavaScript. The V8 API has changed over time and there is no
guarantee it won't again.
Sometime there will be some breaking changes introduced in each new release.

Many application are still running on Node.js v0.10, v0.11, v0.12 is v0.13
There is a real need for some form of abstraction, so that we might
forget a more stable and clean API.

NAN (Native abstraction for Node.js) provides a set of macros and
utilities that achieve two goals:

- 1) Providing a stable API across all Node.js versions
- 2) Providing a set of utilities to help simplify some of the most
difficult parts of Node.js and V8 API's

NAN is not a high level API, but as we have a small utility to simplify

~~asynchronous~~ asynchronous addon development.

~~asynchronous~~ asynchronous addon development.

than module, header files being on the build paths when we compile
our C++ code.

Maybe Type Concept

Maybe Types may or not hold values when we call `ToLocalChecked()` the value will be converted to the real Local

~~NaN starts private~~

NAN doesn't provide a way to cast arrays, probably because the v8 API has never changed about that.

One area where NAN goes a bit beyond simple macros and API redefinition is when dealing with callbacks and asynchronous workers.

Callback object is a NAN class that wraps a standard V8::Function as the most important thing is that it provides ~~function~~ wrapped function from garbage collector.

NAN offers several utility classes that abstract the concept of moving data between the event loop and worker threads. Non-Async Worker

First step is to create a class that inherits from Async Worker and implement

`Execute()` method.

Generally input will be stored on pool for your worker thread.

After the execution that happens on background thread a `HandleOKCallback` or

`HandleErrorCallback` will be executed on the event loop.

NAN provides a utility method called `AsyncQueueWorker` which accepts a `AsyncWorker`

queue and queue's it's execution up such that `Execute` is run only on a worker thread.

This frees us from dealing directly with libuv.

This frees us from dealing directly with libuv.

The asynchronous version of an add-on is far superior in that the heavy CPU computation being done for large N values won't tie up the event loop from working on other things.

AsyncProgressWorker should be a convenient pattern for sending progress updates back to JavaScript during asynchronous execution.

ObjectWraps allow us to pass complete C++ classes to JavaScript. We decorate them with Node/V8 boilerplate code that leaves us susceptible to version issues. ObjectWraps API in fact has been changed through several versions of Node.js -

NAN aims to never fundamentally alter the programming model around C++ code, only to provide a thin layer ~~over~~ around V8 API to ensure backwards and forwards compatibility.

Use NAN wherever possible. It's simpler and safer than trying to keep up with versions of V8 yourself.

Alternatives to Addins

Automation → call your C++ as a standalone app in a child process.

Shared Library → pack your C++ code to a shared library and call these routines directly from javascript

FFI (Foreign Function Interface)

Streaming between Node and C++

Sometimes we use C++ to do some heavy task that actually generates partial results over time.

All partial results are returned immediately through the invoking of a

JS callback

In this case we will use a series of callbacks to build an EventEmitter

Interface -

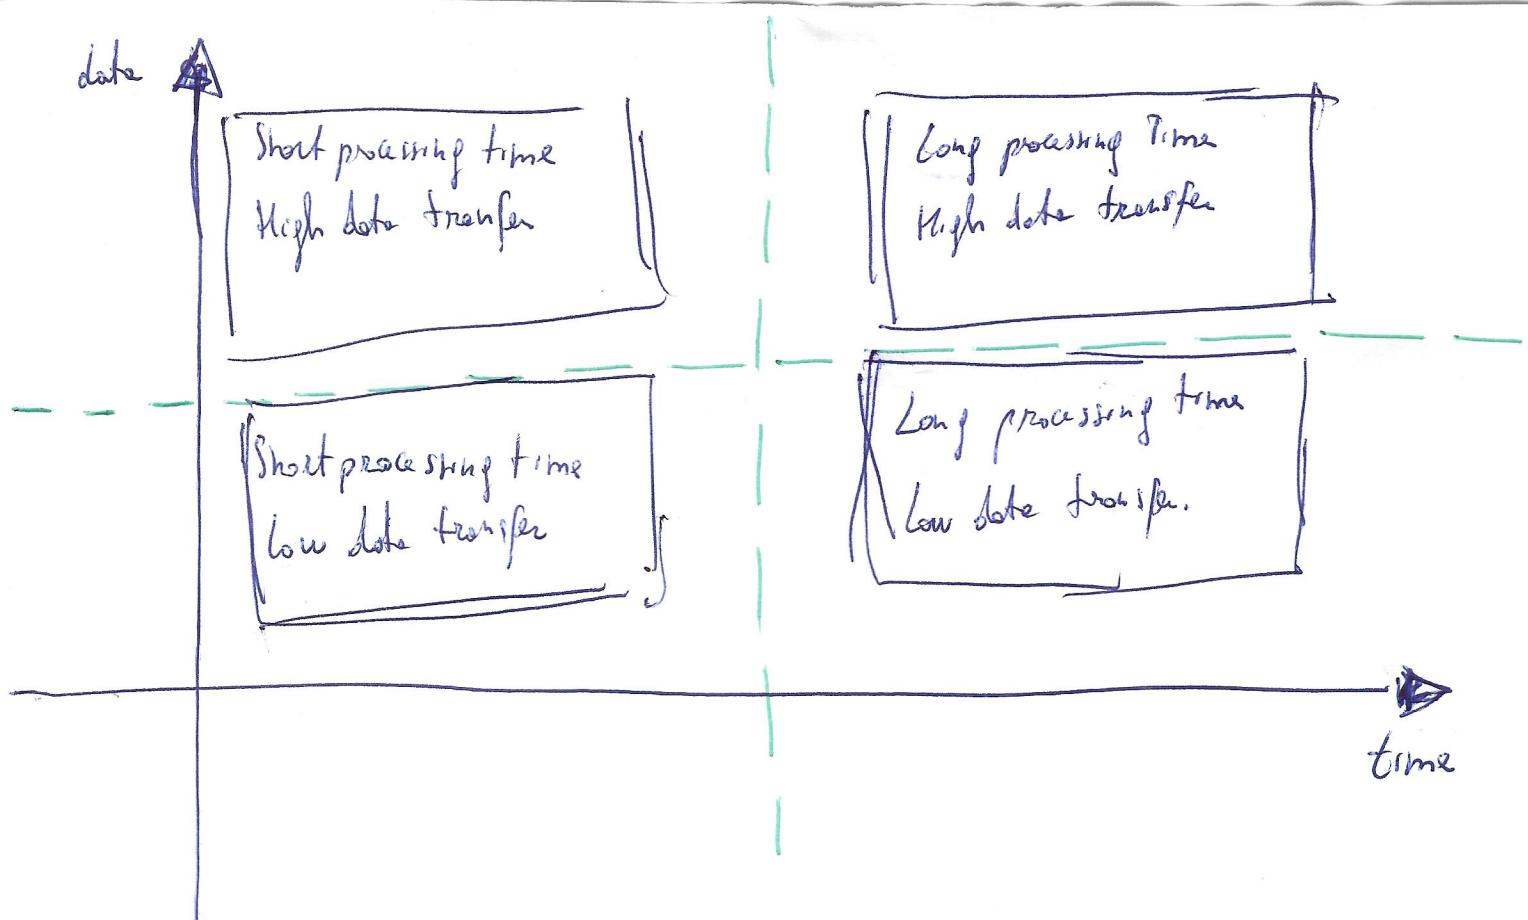
Event - Based Add-Ins

complete progress error

streaming-worker.h

We can use JSON in our Add-Ins

Streaming C++ output



We can categorize our ~~other~~ native Add-ons using two axes

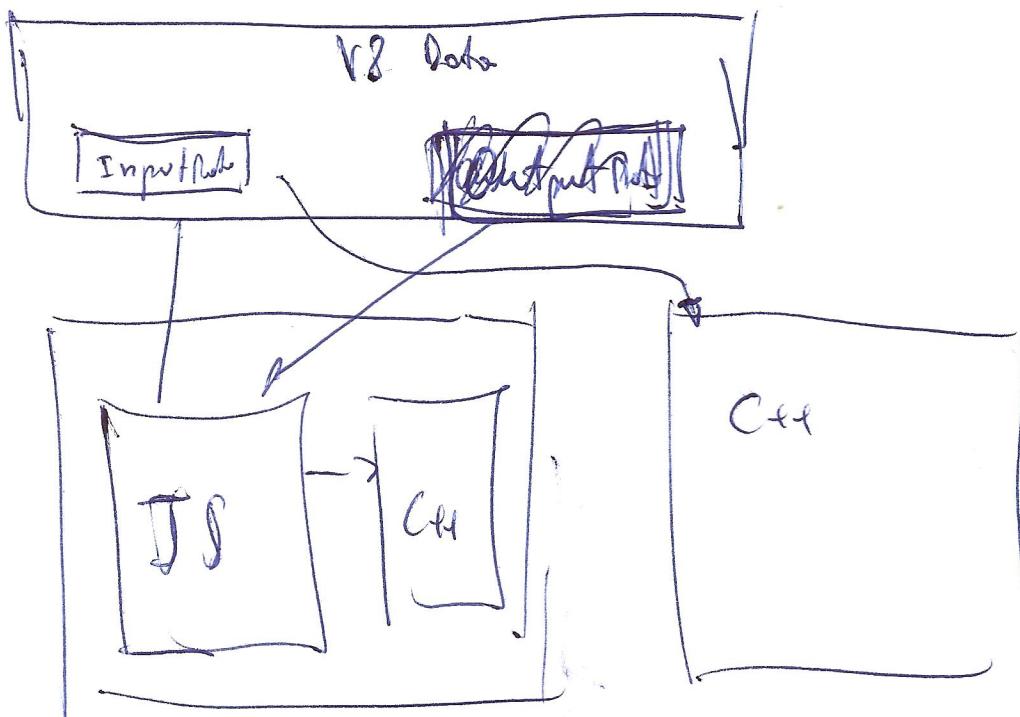
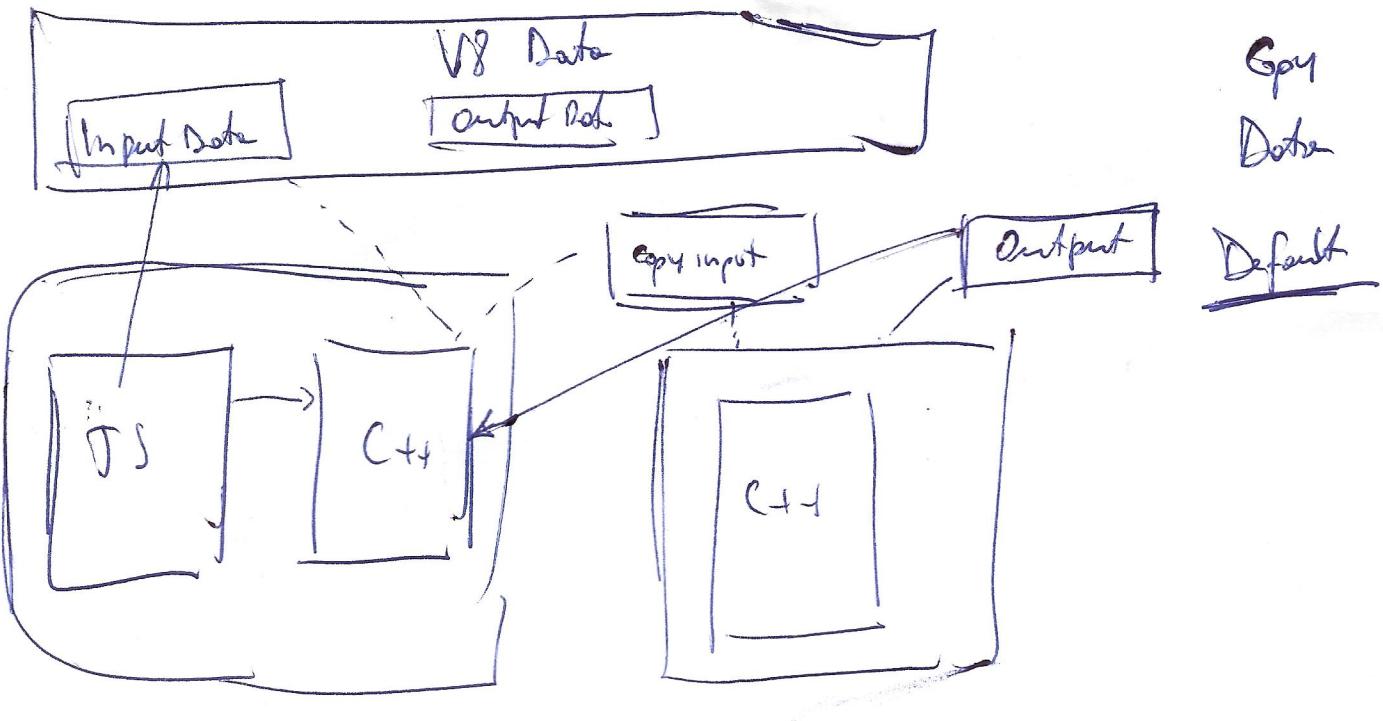
time → Processing time we will spend in the C++ code

data → The amount of data flowing between C++ and JavaScript

Short processing time → your add-on can possibly be synchronous. C++ code running on Worker.js event loop

Long processing time → your add-on ~~can~~ ^{has to be} be asynchronous.

JavaScript primitives are immutable, ~~large~~ cells associated with primitive JavaScript variables cannot be altered by a C++ add-on. The primitive JavaScript ~~variables~~ ^{variables} can be reassigned to new storage cell. Created by C++ but this means changing data always results in new memory allocation.



Instances of the Buffer class are similar to arrays of integers but correspond to fixed-sized, raw memory allocations outside the V8 heap.

You can't ignore the costs of copying data between V8 storage cells and C++ variables. If you aren't careful you can easily kill the performance. ~~cost~~ you might have thought you were getting by dropping into C++ to perform your work.

Buffers offers a way to work with the same data in both JavaScript and C++ thus avoiding the need to create copies.